


Final: AlienGLRenderer: Cloth Simulation





Arrian Chi



Inspiration

AlienGLRenderer

- A sandbox renderer in OpenGL
 - Streamline creation of simple scenes
 - Help me reinforce CG concepts
- Current features:
 - Post-processing
 - Instancing
 - Frustum culling
 - Gltf Model Loading
 - Scene Loading



I occasionally post my progress here: <https://x.com/isacsac2017>


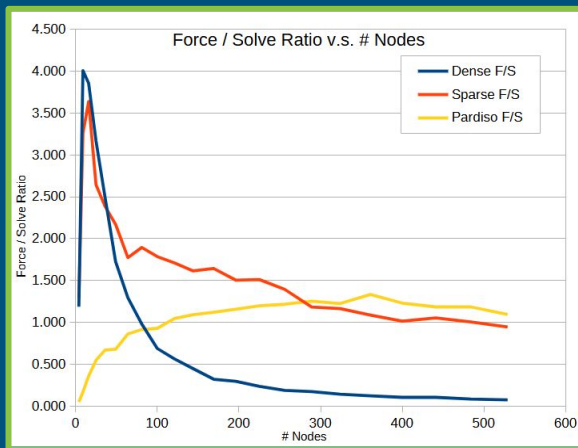

Questions

- How efficiently can we render a cloth with the DEP algorithm?
- How do we implement this?
- What problems arise when implementing the simulation in C++ / OpenGL?

Questions

- How efficiently can we render a cloth with the DEP algorithm?
 - Current Best: 256 - 289 nodes (768 - 867 DOFs) at ~30FPS
 - With a slight caveat maybe more?
- How do we implement this?
 - Use sparse matrices
 - Use faster solvers (ONLY IF THEY ARE THE BOTTLENECK)
- What problems arise when implementing the simulation in C++ / OpenGL?
 - Memory alignment
 - API design issues
 - Numerical Degeneracies

Results - Cont.




Implementation #1

Transcription of python version

- Numpy -> Eigen
- Dense Matrices
- Double -> floats

~30 FPS: 169 - 196 nodes (507 - 588 DOFs)

529 Nodes: 1.145 FPS



Implementation #2

Use Sparse Matrices


- Floats -> Doubles
- Dense Matrices -> Sparse Matrices
- Not Multithreaded
 - Actually didn't help that much
 - Cache coherency (false sharing)
- Fixed a major bug

~30 FPS: 484 nodes (1575 DOFs)

1764 Nodes: 4.098 FPS

0	3	0	0	0
22	0	0	0	17
7	5	0	1	0
0	0	0	0	0
0	0	14	0	8

Values:	22	7	3	5	14	1	17	8
InnerIndices:	1	2	0	2	4	2	1	4
OuterStarts:	0	2	4	5	6	8		



Implementation #3

Third implementation: using PARDISO


- Full Matrix -> Upper triangular
- Column major -> Row major

~30 FPS: 484 nodes (1575 DOFs)

1764 Nodes: 4.435 FPS

	1	2	3	4	5	6	7	8
1	7.		1.			2.	7.	
2		-4.	8.		2.			
3			1.					5.
4				7.			9.	
5					5.	-1.	5.	
6						0.		5.
7							11.	
8								5.



Symmetric Matrix		
IA(k)	JA(K)	A(K)
1	1	7.
5	3	1.
8	6	2.
10	7	7.
12	2	-4.
15	3	8.
17	5	2.
18	3	1.
19	8	5.
	4	7.
	7	9.
	5	5.
	6	-1.
	7	5.
	6	0.
	8	5.
	7	11.
	8	5.



Comments

- OpenGL kinda sucks
 - Memory Alignment
- C++ is hard
 - Template metaprogramming is hard to debug
- Degeneracies
 - Arise when the young modulus is too large, the # nodes too large, the timestep too large etc.
 - Tweaking parameters is needed
- Eigen (reasonably) doesn't have sparse matrix slicing
 - Develop intuition about how matrix slicing works.

Thank you!



References

- [1] D. Baraff and A. Witkin, Large Steps in Cloth Simulation. New York, NY, USA: Association for Computing Machinery, 1 ed., 1998.
- [2] E. Grinspun, A. N. Hirani, M. Desbrun, and P. Schröder, “Discrete shells,” in Proceedings of the 2003 ACM SIGGRAPH/Eurographics Symposium on Computer Animation, SCA '03, (Goslar, DEU), p. 62–67, Eurographics Association, 2003.
- [3] R. Tamstorf and E. Grinspun, “Discrete bending forces and their jacobians,” Graph. Models, vol. 75, p. 362–370, Nov. 2013.
- [4] J. Weil, “The synthesis of cloth objects,” ACM Siggraph Computer Graphics, vol. 20, no. 4, pp. 49–54, 1986.
- [5] X. Provot et al., “Deformation constraints in a mass-spring model to describe rigid cloth behaviour,” 1995.
- [6] D. Terzopoulos, J. Platt, A. Barr, and K. Fleischer, “Elastically deformable models,” SIGGRAPH Comput. Graph., vol. 21, p. 205–214, Aug. 1987.
- [7] S. MK Jawed, Lim, “Discrete simulation of slender structures,” 2023.

References

- [8] G. Guennebaud, B. Jacob, et al., Eigen: A C++ template library for linear algebra. Eigen Contributors, 2024. Accessed: 2024-11-20.
- [9] L. K. Contributors, False Sharing in the Linux Kernel, 2024. Accessed: 2024-11-20.
- [10] D. Pasadakis, M. Bollhöfer, and O. Schenk, “Sparse quadratic approximation for graph learning,” IEEE Transactions on Pattern Analysis and Machine Intelligence, vol. 45, no. 9, pp. 11256–11269, 2023.
- [11] A. Eftekhari, D. Pasadakis, M. Bollhöfer, S. Scheidegger, and O. Schenk, “Block-enhanced precision matrix estimation for large-scale datasets,” Journal of Computational Science, vol. 53, p. 101389, 2021.
- [12] L. Gaedke-Merzhäuser, J. van Nieuwerkerk, O. Schenk, and H. Rue, “Parallelized integrated nested laplace approximations for fast bayesian inference,” 2022.
- [13] C. Caruth, “There are no zero-cost abstractions.” CPPCon 2019, 2019.
- [14] Khronos Group, OpenGL Documentation. Khronos Group. Accessed: 2024-11-20.